

```
# Import libraries import re for text cleaning (regex) import nltk for
tokenization and stopwords from collections import Counter, defaultdict #
for counting n -grams import math for probability and perplexity calculations
```

Explanation :

re remove punctuation/ numbers nltk tokenize words, handle stopwords collections. Counter count word frequencies collections. defaultdict store conditional probabilities math log probabilities, perplexity calculations

```
# Example dataset: a text corpus (at
with , "r") as f: text = f.read()
# The file 'dataset.txt' was not fou
# To fix this, we will temporarily u
# Please replace this placeholder wi #
or upload 'dataset.txt' to your Cc
```

text -
This is a placeholder text for your You
should replace this with a much to
properly test the language model. an
article, or any other long piece If you
have a 'dataset.txt' file, pl and then
uncomment the original 'wit

A language model is a probability di
Given such a sequence, say of length
Language models are useful for a var
information retrieval, and many natu
The simplest type of language model
An n-gram is a contiguous sequence c
For example, in the phrase "the quic and
"the quick brown" is a trigram (

N-gram models estimate the probabili
This is done by counting how often s For
example, to calculate P(word_k | I we count
the occurrences of the sequ and divide it by
the count of the se This approach relies on
the Markov a depends only on the previous
n-l wor

Smoothing techniques are often used
Common smoothing methods include Add
Good-Turing smoothing, and Kneser-
Ne These methods redistribute some
prob preventing zero probabilities that

```
# If you have 'dataset.txt' uploaded
# with open("dataset.txt", "r") as f
text = f.read()

# Clean unnecessary lines
text = text.strip()
```

```
# Display sample
print(text [: 500])
```

This is a placeholder text for your dataset.

You should replace this with a much larger text corpus, ideally at least 1500 words, to properly test the language model. For example, you can paste text from a book, an article, or any other long piece of writing here.

If you have a 'dataset.txt' file, please upload it to your Colab environment, and then uncomment the original 'with open("dataset.txt", "r") as f: text = f.read()' line.

A language model is a probability distribution over a sequence of

+ Gemini

```
import nltk
nltk.download('punkt')
nltk.download('stopwords')
nltk.download('punkt_tab') # Add this

def preprocess(text):
    Lowercase_text =
        text.lower()
    # Remove punctuation and
    numbers_text -- re.sub(r'[^\w\s]', ''
    # Tokenize tokens =
        nltk.
        word_tokenize(text) # Optionally
        remove_stopwords_stopwords =
        set(nltk.corpus.stopwords)
        tokens = [t
        for t in tokens if t # Add start/end
        tokens processed for sentence in
        nltk.sent_tokenize(words =
        nltk.word_tokenize(s processed.
        append(["<s>"]) + return processed
```

processed_text = preprocess(text)

[nltk_data] Downloading package punkt to /root/nltk_data..

[nltk_data] Package punkt is already up-to-date!

[nltk_data] Downloading package stopwords to
/root/nltk_data.

[nltk_data] Package stopwords is already up-to-date!

[nltk_data] Downloading package punkt_tab to
/root/nltk_data. [nltk_data] Unzipping
tokenizers/punkt_tab.zip.

```

def build_ngram(tokens, n):
    ngrams = []
    for sentence in tokens:
        for i in range(len(sentence)-n+1):
            ngrams.append(tuple(sentence[i:i+n]))
    return Counter(ngrams)

# Unigram, Bigram, Trigram
unigrams = build_ngram(processed_text, 1)
bigrams = build_ngram(processed_text, 2)
trigrams = build_ngram(processed_text, 3)

print("Sample Unigrams:", list(unigrams.items())[:10])
print("Sample Bigrams:", list(bigrams.items())[:10])

```

Sample Unigrams: [('<s>',), 1], ('this',), 4), ('is',), 7), ('a',), 16), ('placeholder',), 1)
Sample Bigrams: [('<s>', 'this'), 1], ('this', 'is'), 2), ('is', 'a'), 5), ('a', 'placeholder')

b.research.google.com/drive/1EdT24GZ8twoyudV9EZcLPz3UirfG6Gbr#scrollTo=T3dPUyXYMLRS&printMode=true

```

def laplace_smoothing(count, total, vocab_size) : return
    (count + 1) / (total + vocab_size)

```

```

def sentence_probability(sentence, model, n, vocab_size) :
    tokens = lower() + prob = 1.0 for i in range(len(tokens)-n+1) : ngram = tokens[i:i+n] count = model[ngram] total = sum(model.values()) prob *= laplace_smoothing(count, total, vocab_size) return prob

```

sentences

```

"The cat sat on the mat." ,
"Students are learning language models." ,
"Artificial intelligence is powerful." ,
"Dogs bark loudly at night." ,
"The quick brown fox jumps over the lazy dog."

```

```

for s in sentences: print ("Sentence: " + s) print( "Unigram Probability: "
    sentence_probability(s, unigrams, 1, len(unigrams))) print( "Bigram
Probability: " + sentence_probability(s, bigrams, 2, len(bigrams))) print(
    "Trigram Probability: " + sentence_probability(s, trigrams, 3, len(trigrams)))
print()

```

Sentence: The cat sat on the mat.

Unigram Probability: 5.944485537594832e-21 Bigram

Probability: 2.6875722256150096+22

Trigram Probability: 3.109854318894012e-20

Sentence: Students are learning language models.

Unigram Probability: 7.925571316962695e-20 Bigram

Probability: 7.941775926692354e-20

Trigram Probability: 1.903230843163135+17

Sentence: Artificial intelligence is powerful.

Unigram Probability: 5.612478651123212e-18

Bigram Probability: 2.346794786337591e-17

Trigram Probability: 1.1647772760158386e-14

Sentence: Dogs bark loudly at night. Unigram
 Probability: 2 .935396784060257e- 21 Bigram
 Probability: 3 .970887963346177e- 20
 Trigram Probability: 1.903230843163135e-17

Sentence: The quick brown fox jumps over the lazy dog.
 Unigram Probability: 6.531483454005439e-28 Bigram
 Probability: 7 .811761079006753e- 30
 Trigram Probability: 8 .140236616337258+28

```
def perplexity(sentence, model, n, vocab_size):
    tokens = + nltk.word_tokenize(sentence.lower()) + log_
    prob for i in range(len(tokens)-n+1) : ngram =      i+n]  )
    count - model [ngram] total = sum(model.values())
```

3/4

```
    prob = laplace_smoothing(count, total, vocab_size)
    log_prob += math.log(prob)
    return math.exp(-log_prob/len(tokens))

for s in sentences:
    print("Sentence:", s)
    print("Unigram Perplexity:", perplexity(s, unigrams, 1, len(unigrams)))
    print("Bigram Perplexity:", perplexity(s, bigrams, 2, len(bigrams)))
    print("Trigram Perplexity:", perplexity(s, trigrams, 3, len(trigrams)))
    print()
```

Sentence: The cat sat on the mat.
 Unigram Perplexity: 176.73421536811202
 Bigram Perplexity: 249.3088880735985
 Trigram Perplexity: 147.05281019460824

Sentence: Students are learning language models.
 Unigram Perplexity: 244.13001494881965
 Bigram Perplexity: 244.06769312592664
 Trigram Perplexity: 123.04488553884742

Sentence: Artificial intelligence is powerful.
 Unigram Perplexity: 291.3442633548697
 Bigram Perplexity: 237.49059315444097
 Trigram Perplexity: 97.8445702538009

Sentence: Dogs bark loudly at night.
 Unigram Perplexity: 368.58838727249776
 Bigram Perplexity: 266.15770664759424
 Trigram Perplexity: 123.04488553884742

Sentence: The quick brown fox jumps over the lazy dog.
 Unigram Perplexity: 184.2534721981669
 Bigram Perplexity: 266.4429626201017
 Trigram Perplexity: 180.90347703814962

